

Class09: Candy Mini-Project

Kenny Dang

Table of contents

Exploring the correlation structure	13
Principal Component Analysis	16
The loadings plot for PC1	18

```
candy = read.csv("candy-data.csv", row.names= 1)  
  
head(candy)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crisp	pedric	icewafer
100 Grand	1	0	1			0	0		1
3 Musketeers	1	0	0			0	1		0
One dime	0	0	0			0	0		0
One quarter	0	0	0			0	0		0
Air Heads	0	1	0			0	0		0
Almond Joy	1	0	0			1	0		0
	hard	bar	pluribus	sugarpercent	pricepercent	percent	winpercent		
100 Grand	0	1	0	0.732	0.860	66.97173			
3 Musketeers	0	1	0	0.604	0.511	67.60294			
One dime	0	0	0	0.011	0.116	32.26109			
One quarter	0	0	0	0.011	0.511	46.11650			
Air Heads	0	0	0	0.906	0.511	52.34146			
Almond Joy	0	1	0	0.465	0.767	50.34755			

Q1. How many different candy types are in this dataset?

There are 85 different candy types in this dataset.

```
nrow(candy)
```

```
[1] 85
```

There are 85 rows in this dataset.

Q2. How many fruity candy types are in the dataset?

There are 38 fruit candy types in the dataset.

```
sum(candy$fruity)
```

```
[1] 38
```

Q3. what is your favorite candy (other than Twix) in the dataset and what is it's `winpercent` value?

The `winpercent` value for my favorite candy (Warheads) is 39.0119.

```
candy["Warheads", ]$winpercent
```

```
[1] 39.0119
```

Q4. What is the `winpercent` value for “Kit Kat”?

The `winpercent` value for “Kit Kats” is 76.7686.

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

Q5. What is the `winpercent` value for “Tootsie Roll Snack Bars”?

The `winpercent` value for “Tootsie Roll Snack Bars” is 49.6535.

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

```
library("skimr")
```

```
Warning: package 'skimr' was built under R version 4.4.3
```

```
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

winpercent is the variable that appears to be on a different scale as the values appear to be between 0-100. All the other variables appear to have values that are between 0-1.

```
summary(candy)
```

	chocolate	fruity	caramel	peanutyalmond
Min.	:0.0000	:0.0000	:0.0000	:0.0000
1st Qu.	:0.0000	:0.0000	:0.0000	:0.0000
Median	:0.0000	:0.0000	:0.0000	:0.0000
Mean	:0.4353	:0.4471	:0.1647	:0.1647
3rd Qu.	:1.0000	:1.0000	:0.0000	:0.0000
Max.	:1.0000	:1.0000	:1.0000	:1.0000
	nougat	crispedrice	wafers	hard
Min.	:0.00000	:0.00000	:0.0000	:0.0000
1st Qu.	:0.00000	:0.00000	:0.0000	:0.0000
Median	:0.00000	:0.00000	:0.0000	:0.0000
Mean	:0.08235	:0.08235	:0.1765	:0.2471
3rd Qu.	:0.00000	:0.00000	:0.0000	:0.0000
Max.	:1.00000	:1.00000	:1.0000	:1.0000
	pluribus	sugarpercent	pricepercent	winpercent
Min.	:0.0000	:0.0110	:0.0110	:22.45
1st Qu.	:0.0000	:0.2200	:0.2550	:39.14
Median	:1.0000	:0.4650	:0.4650	:47.83
Mean	:0.5176	:0.4786	:0.4689	:50.32
3rd Qu.	:1.0000	:0.7320	:0.6510	:59.86
Max.	:1.0000	:0.9880	:0.9760	:84.18

Q7. What do you think a zero and one represent for the “candy\$chocolate” column?

1 means the candy contains chocolate while 0 means the candy does not contain chocolate.

```
candy$chocolate
```

```
[1] 1 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 0 1 1 0 0 0 1 1 0 1 1 1  
[39] 1 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 0 1 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 0 1 1  
[77] 1 1 0 1 0 0 0 0 1
```

```
##Exploratory Analysis
```

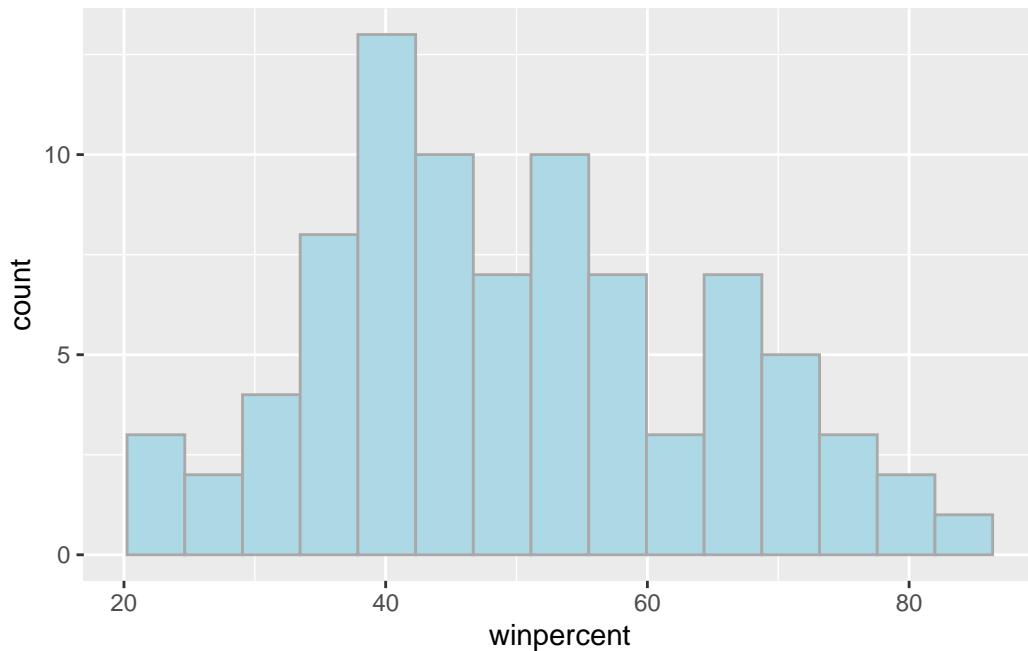
Q8. Plot a histogram of winpercent values using both base R an ggplot2.

Plot is shown below.

```
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.4.3

```
ggplot(candy) +  
  aes(winpercent) +  
  geom_histogram(bins=15, fill="lightblue", col="darkgray")
```



Q9. Is the distribution of winpercent values symmetrical?

Since the mean is not equal to the median, this means that the winpercent values are not symmetrical. And it can also be seen on the graph in question 9 that the winpercent values are not symmetrical.

```
mean(candy$winpercent)
```

[1] 50.31676

```
median(candy$winpercent)
```

[1] 47.82975

The distribution is not symmetrical.

Q10. Is the center of the distribution above or below 50%?

The center of distribution is below 50%.

```
mean(candy$winpercent)
```

```
[1] 50.31676
```

```
median(candy$winpercent)
```

```
[1] 47.82975
```

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

On average, chocolate candy is higher than the fruit candy.

Steps to solve this: 1. Find all chocolate candy in the dataset 2. Extract or find their winpercent values 3. Calculate the mean of these values

4. Find all fruit candy
5. Find their winpercent values
6. Calculate their mean value

```
chocolate_candy <- candy[candy$chocolate == 1,]
chocolate_winpercent <- chocolate_candy$winpercent
mean_chocolate <- mean(chocolate_winpercent, na.rm = TRUE)
mean_chocolate
```

```
[1] 60.92153
```

```
fruit_candy <- candy[candy$fruit == 1,]
fruit_winpercent <- fruit_candy$winpercent
mean_fruit <- mean(fruit_winpercent, na.rm = TRUE)
mean_fruit
```

```
[1] 44.11974
```

Q12. Is this difference statistically significant?

The difference is statistically significant as the p-value is $2.871*10e-08$ which is less than 0.05.

```
t.test(chocolate_winpercent, fruit_winpercent)
```

Welch Two Sample t-test

```
data: chocolate_winpercent and fruit_winpercent
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Q13. What are the five least liked candy types in this set?

The five least liked candy types in this dataset are Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, and Jawbusters.

```
inds <- order(candy$winpercent)
candy[inds,] [1:5,]
```

	chocolate	fruity	caramel	peanut	yalmond	nougat		
Nik L Nip	0	1	0		0	0		
Boston Baked Beans	0	0	0		1	0		
Chiclets	0	1	0		0	0		
Super Bubble	0	1	0		0	0		
Jawbusters	0	1	0		0	0		
	crisped	rice	wafer	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0		1		0.197	0.976
Boston Baked Beans	0	0	0		1		0.313	0.511
Chiclets	0	0	0		1		0.046	0.325
Super Bubble	0	0	0		0		0.162	0.116
Jawbusters	0	1	0		1		0.093	0.511
	winpercent							
Nik L Nip	22.44534							
Boston Baked Beans	23.41782							
Chiclets	24.52499							
Super Bubble	27.30386							
Jawbusters	28.12744							

Q14. What are the top 5 all time favorite candy types out of this set?

The top 5 all time favorite candy types of this data set are: snickers, kit kat, twix, reese's minatures, reese's peanut butter cup.

```
inds <- order(candy$winpercent, decreasing = TRUE)
candy[inds, ][1:5, ]
```

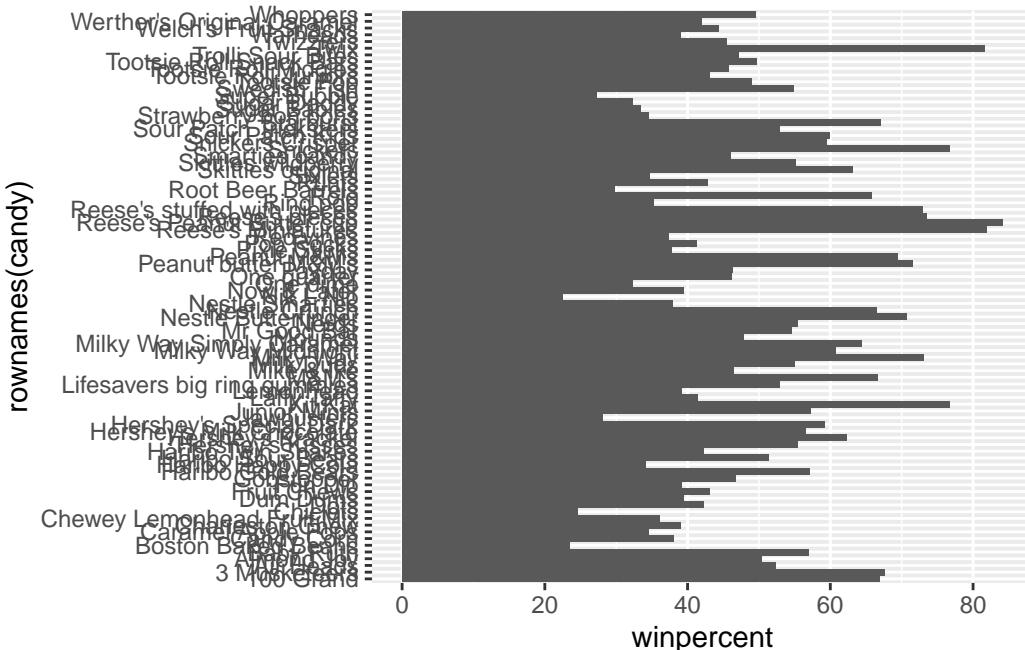
	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1
	crispedrice	wafers	hard bar	pluribus	sugar	percent
Reese's Peanut Butter cup	0	0	0	0	0.720	
Reese's Miniatures	0	0	0	0	0.034	
Twix	1	0	1	0	0.546	
Kit Kat	1	0	1	0	0.313	
Snickers	0	0	1	0	0.546	
	price	percent	winpercent			
Reese's Peanut Butter cup	0.651	84.18029				
Reese's Miniatures	0.279	81.86626				
Twix	0.906	81.64291				
Kit Kat	0.511	76.76860				
Snickers	0.651	76.67378				

Q15. Make a first barplot of candy ranking based on winpercent values.

barplot shown below.

```
library (ggplot2)

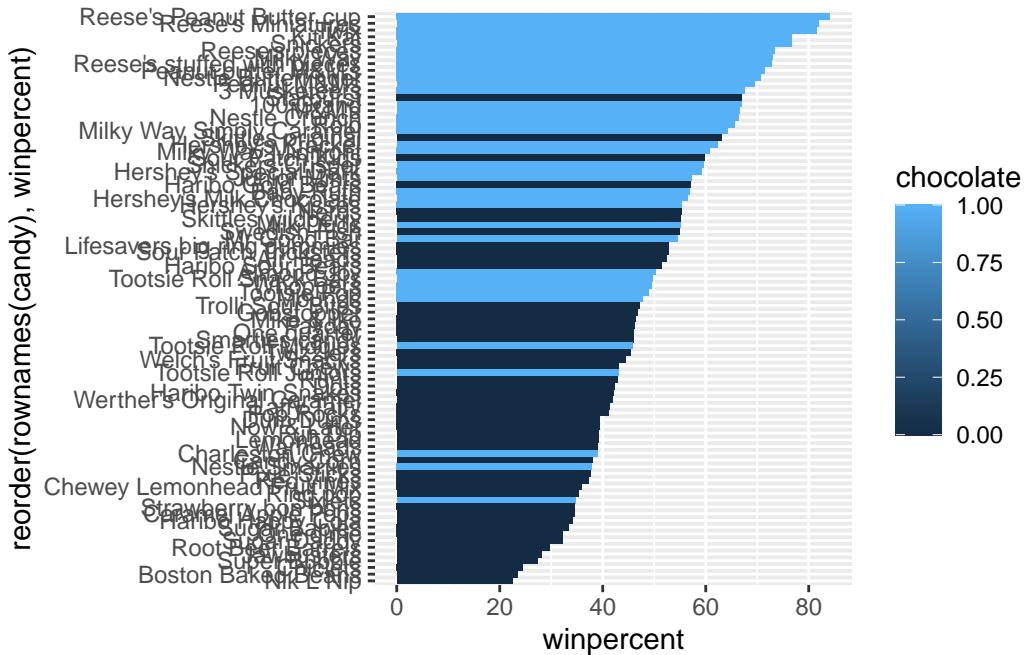
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

`reorder()` function used below.

```
ggplot(candy) +  
  aes(winpercent,  
      reorder(rownames(candy), winpercent),  
      fill = chocolate) +  
  geom_col()
```



![Figure caption: my second barplot] (barplot2.png)

I want custom colors that I pick so we need to make this ourselves...

```
my_cols <- rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] <- "brown"
my_cols[as.logical(candy$fruity)] <- "pink"
my_cols[10] <- "red"
my_cols
```

```
[1] "brown"      "brown"       "black"        "black"        "pink"        "brown"
[7] "brown"      "black"       "black"        "red"          "brown"       "pink"
[13] "pink"        "pink"        "pink"         "pink"         "pink"        "pink"
[19] "pink"        "black"       "pink"         "pink"         "chocolate"   "brown"
[25] "brown"       "brown"       "pink"         "chocolate"   "brown"       "pink"
[31] "pink"        "pink"        "chocolate"   "chocolate"   "pink"        "chocolate"
[37] "brown"       "brown"       "brown"        "brown"        "brown"       "pink"
[43] "brown"       "brown"       "pink"         "pink"         "brown"       "chocolate"
[49] "black"        "pink"        "pink"         "chocolate"   "chocolate"   "chocolate"
[55] "chocolate"   "pink"        "chocolate"   "black"        "pink"        "chocolate"
[61] "pink"         "pink"        "chocolate"   "pink"         "brown"       "brown"
[67] "pink"         "pink"        "pink"         "pink"         "black"       "black"
```

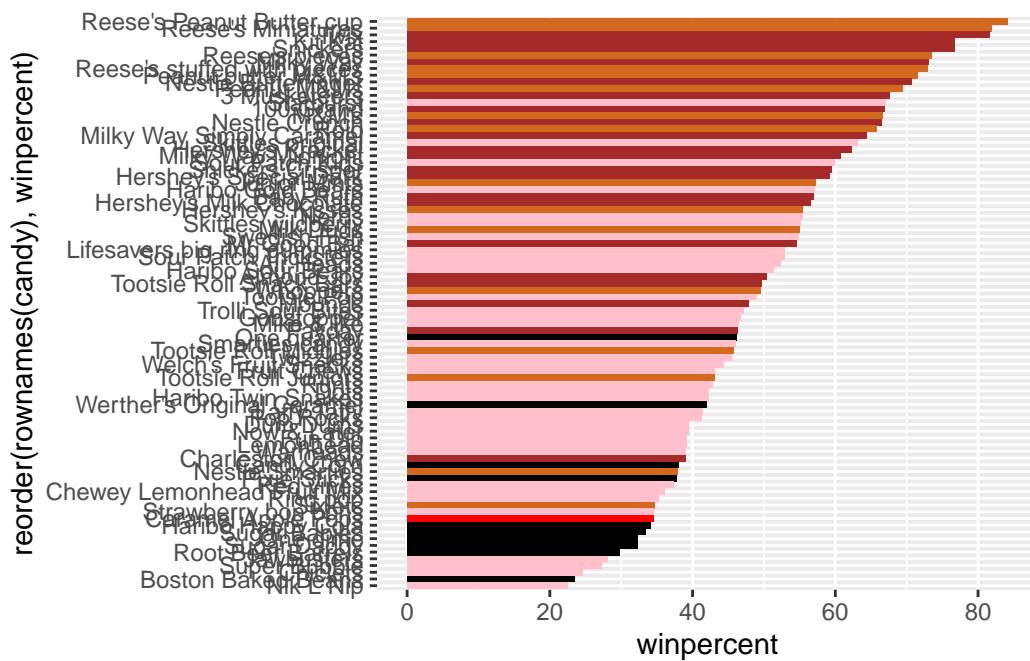
```
[73] "pink"      "pink"      "pink"      "chocolate" "chocolate" "brown"
[79] "pink"      "brown"     "pink"      "pink"       "pink"      "black"
[85] "chocolate"
```

```
my_cols <- rep("black", nrow(candy)) my_cols[candy$chocolate == 1] <- "chocolate" my_cols[candy$bar == 1] <- "brown" my_cols[candy$fruity] <- "blue" my_cols[10] <- "red" my_cols
```

Q17. What is the worst ranked chocolate candy?

The worst ranked chocolate candy is Sixlets.

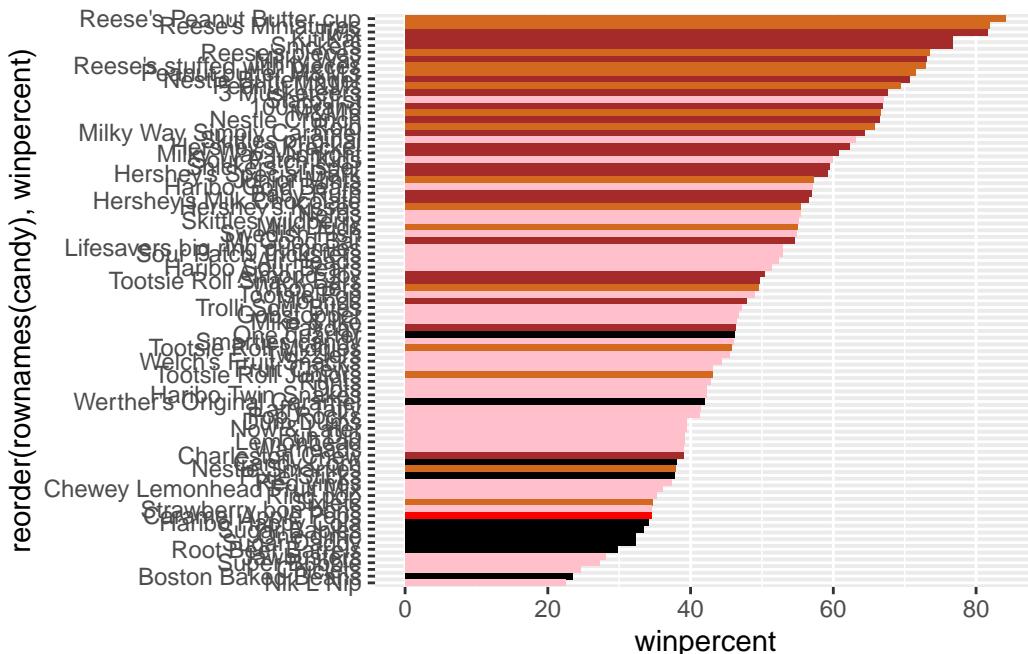
```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill= my_cols)
```



Q18. What is the best ranked fruity candy?

The best ranked fruity candy is Starburst.

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill= my_cols)
```



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

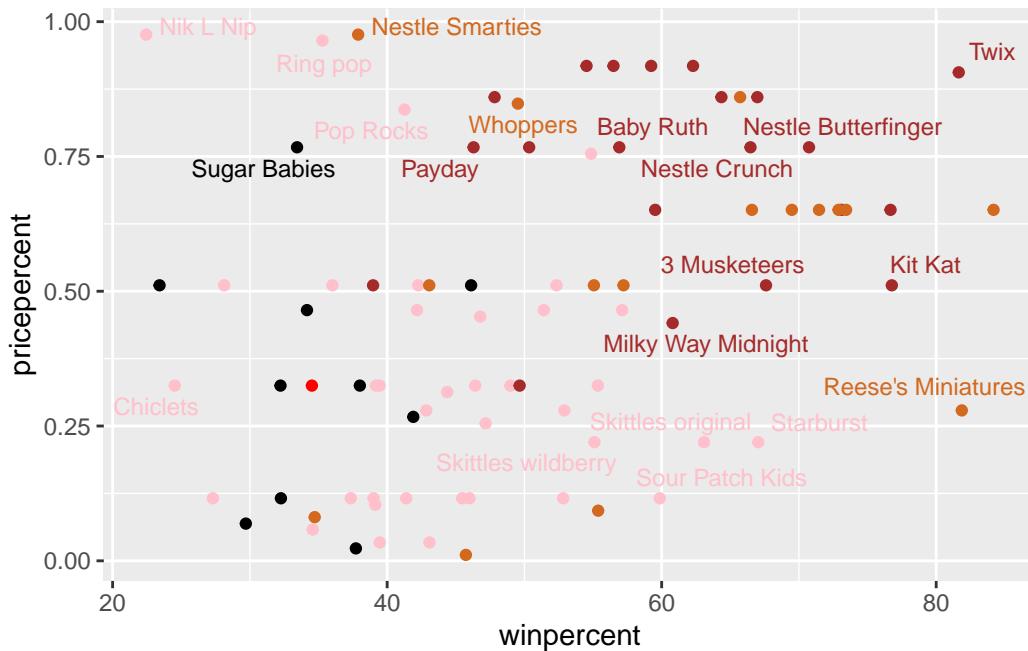
Reese's Miniatures, which is a chocolate candy, is the highest ranked in terms of winpercent for the least money.

```
library(ggrepel)
```

Warning: package 'ggrepel' was built under R version 4.4.3

```
# How about a plot of win vs price
ggplot(candy) +
  aes(x = winpercent, y=pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

The top 5 most expensive candy types in the dataset are Nik L Nip, Nestle Smarties, Ring pop, and Hershey's Krackel, and Hershey's Milk Chocolate.

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head( candy[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Exploring the correlation structure

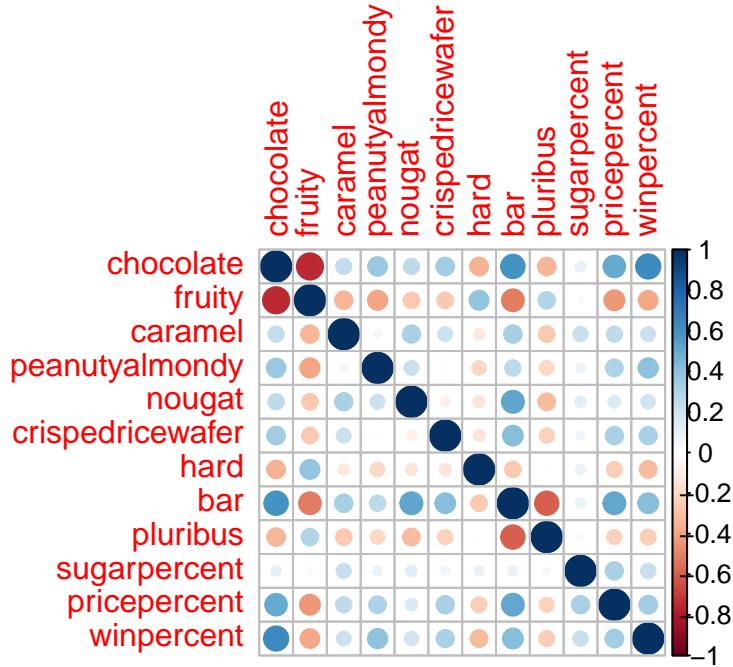
Pearson correlation values range from -1 to +1

```
library (corrplot)
```

```
Warning: package 'corrplot' was built under R version 4.4.3
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```

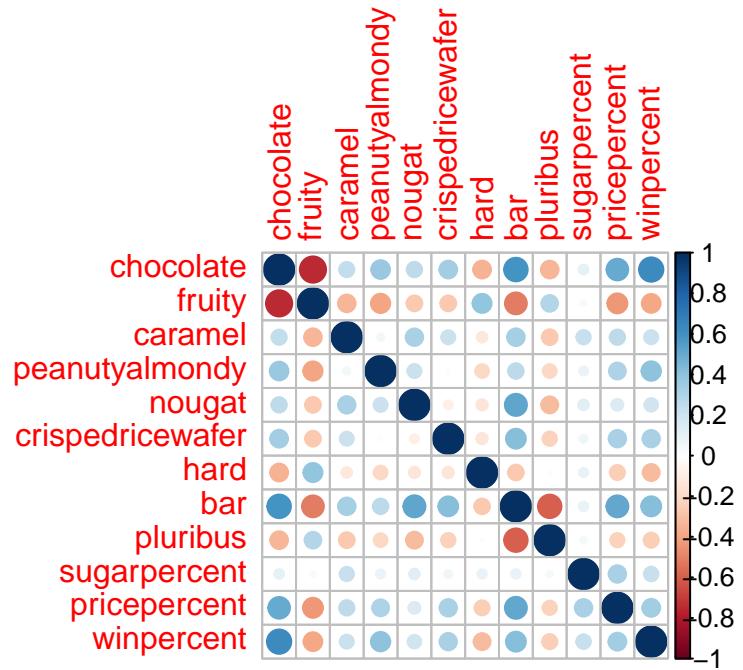


Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)

The two variables that are anti-correlated are chocolate and fruity as their relationship is represented by a red circle (negative value).

```
library (corrplot)

cij <- cor(candy)
corrplot(cij)
```

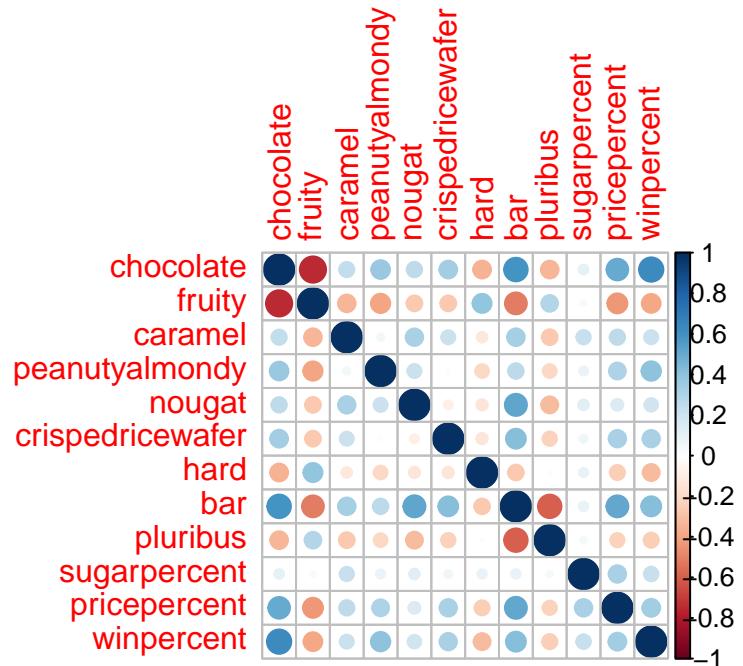


Q23. Similarly, what two variables are most positively correlated?

The two variables that are most positively correlated are chocolate and winpercent as their relationship is represented by a blue circle (positive value).

```
library (corrplot)

cij <- cor(candy)
corrplot(cij)
```



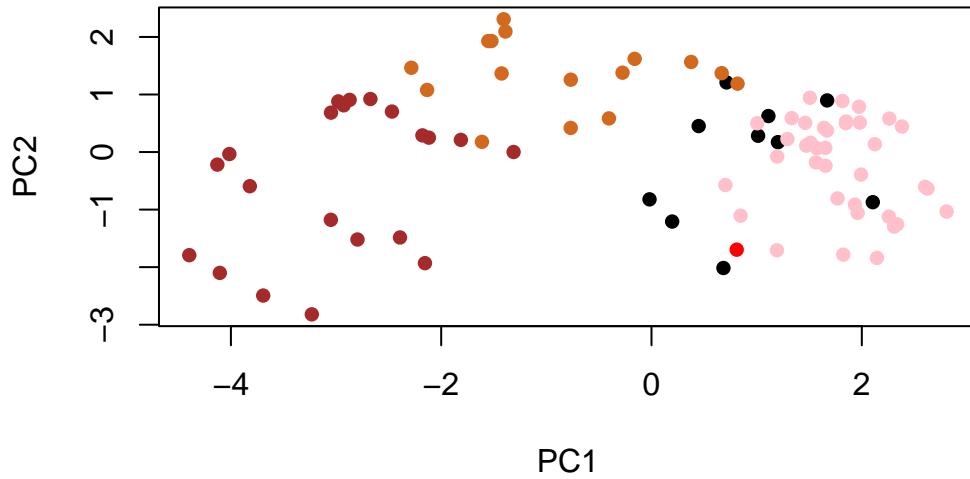
Principal Component Analysis

```
pca <- prcomp(candy, scale=T)
summary (pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

```
plot (pca$x[,1:2], col=my_cols, pch = 16)
```

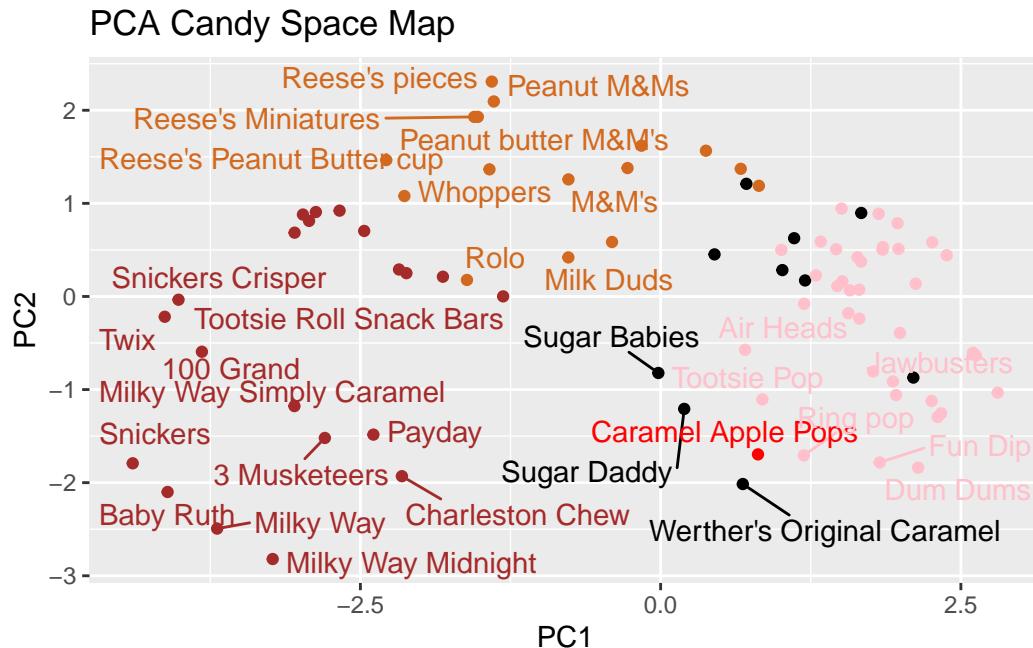


The main results figure: the PCA score plot:

```
library (ggrepel)

ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col = my_cols) +
  geom_text_repel(col = my_cols) +
  labs (title = "PCA Candy Space Map")
```

Warning: ggrepel: 54 unlabeled data points (too many overlaps). Consider increasing max.overlaps

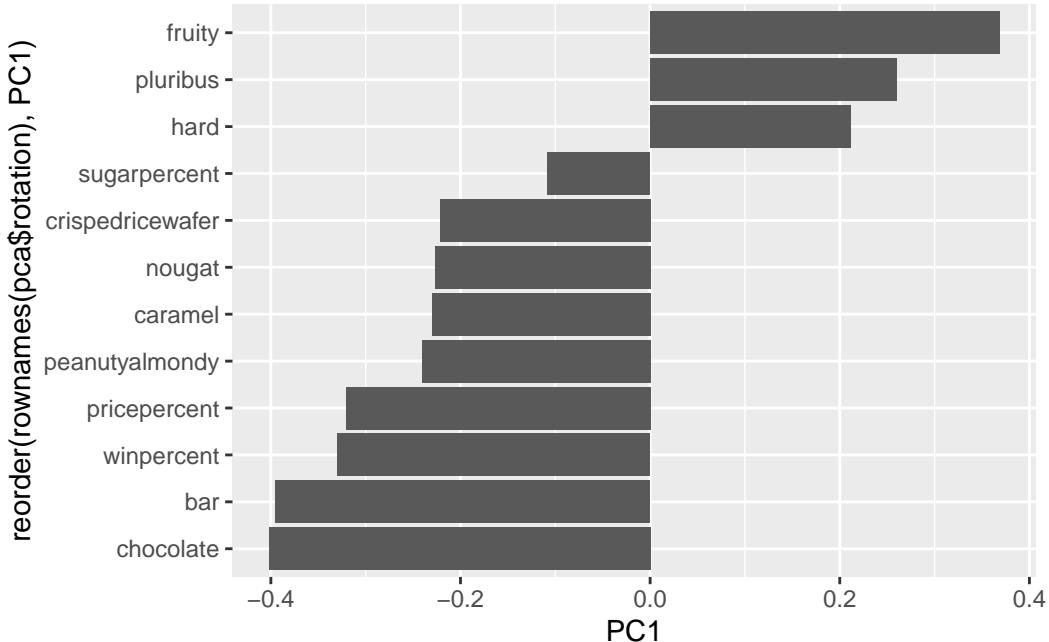


Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

The variables that are picked up strongly by PC1 in the positive direction are “fruity”, “pluribus”, and “hard”. Yes, this makes sense. This relationship was highlighted earlier when creating the correlation structure which essentially depicted whether or not two variables had a positive correlation or a negative correlation. This can also be seen on the PCA scores plot as well.

The loadings plot for PC1

```
ggplot(pca$rotation) +
  aes(PC1, reorder(rownames(pca$rotation), PC1)) +
  geom_col()
```



Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

The characteristics that appear to make a “winning” candy are that it contains chocolate and it is a bar-style candy. The candy also being priced on the lower end (pricepercent) makes it more popular among consumers. These characteristics make a “winning” candy which beats other types of candy such as fruity and hard candies. The scatterplot comparing winpercent vs pricerpercent shows a cluster of chocolate bar style candies clustering towards the high end of the winpercent axis. The correlation structure shows that the variables chocolate and bar have a positive correlation and the variables chocolate and winpercent have a positive correlation. The PCA bar plot also shows the strong positive relationship between the variables chocolate, bar, and winpercent.