

# class010: Structural Bioinformatics (pt. 1)

Kenny Dang (PID: A18544481)

## The PDB database

The [Protein Data Bank \(PDB\)](#) is the main repository of biomolecular structure data. Let's see what is in it:

```
PDB <- read.csv("pdb_stats.csv", row.names = 1)
PDB
```

	X.ray	EM	NMR	Integrative	Multiple.methods	Neutron
Protein (only)	178795	21825	12773	343	226	84
Protein/Oligosaccharide	10363	3564	34	8	11	1
Protein/NA	9106	6335	287	24	7	0
Nucleic acid (only)	3132	221	1566	3	15	3
Other	175	25	33	4	0	0
Oligosaccharide (only)	11	0	6	0	1	0
	Other	Total				
Protein (only)	32	214078				
Protein/Oligosaccharide	0	13981				
Protein/NA	0	15759				
Nucleic acid (only)	1	4941				
Other	0	237				
Oligosaccharide (only)	4	22				

```
head(PDB)
```

	X.ray	EM	NMR	Integrative	Multiple.methods	Neutron
Protein (only)	178795	21825	12773	343	226	84
Protein/Oligosaccharide	10363	3564	34	8	11	1
Protein/NA	9106	6335	287	24	7	0
Nucleic acid (only)	3132	221	1566	3	15	3
Other	175	25	33	4	0	0

Oligosaccharide (only)	11	0	6	0	1	0
Other		Total				
Protein (only)	32	214078				
Protein/Oligosaccharide	0	13981				
Protein/NA	0	15759				
Nucleic acid (only)	1	4941				
Other	0	237				
Oligosaccharide (only)	4	22				

Q1. What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy.

81% of structures are solved by X-ray. 13% of structures are solved by Electron Microscopy.

```
n.sums <- colSums(PDB)
n <- n.sums/n.sums["Total"]
round(n, digits = 2)
```

	X.ray	EM	NMR	Integrative
	0.81	0.13	0.06	0.00
Multiple.methods		Neutron	Other	Total
	0.00	0.00	0.00	1.00

Q2. What proportion of structures in the PDB are protein?

The proportion of structures in the PDB that are protein are 86%.

```
protein_total <- PDB["Protein (only)", "Total"]
protein_total
```

```
[1] 214078
```

```
overall_total <- sum(PDB$Total)
overall_total
```

```
[1] 249018
```

```
proportion <- protein_total / overall_total
proportion
```

```
[1] 0.8596889
```

```
round(proportion, digits = 2)
```

```
[1] 0.86
```

Q3. Type HIV in the PDB website search box on the home page and determine how many HIV-1 protease structures are in the current PDB?

After typing HIV in the PDB website search box, there are 2,427 HIV-1 protease structures in the current PDB.

What is the total number of entries in the PDB

The total number of entries in the PDB are 249018.

```
n.sums["Total"]
```

```
Total  
249018
```

## Using Molstar

We can use the main [Molstar viewer online] ( <https://molstar.org/viewer/>)

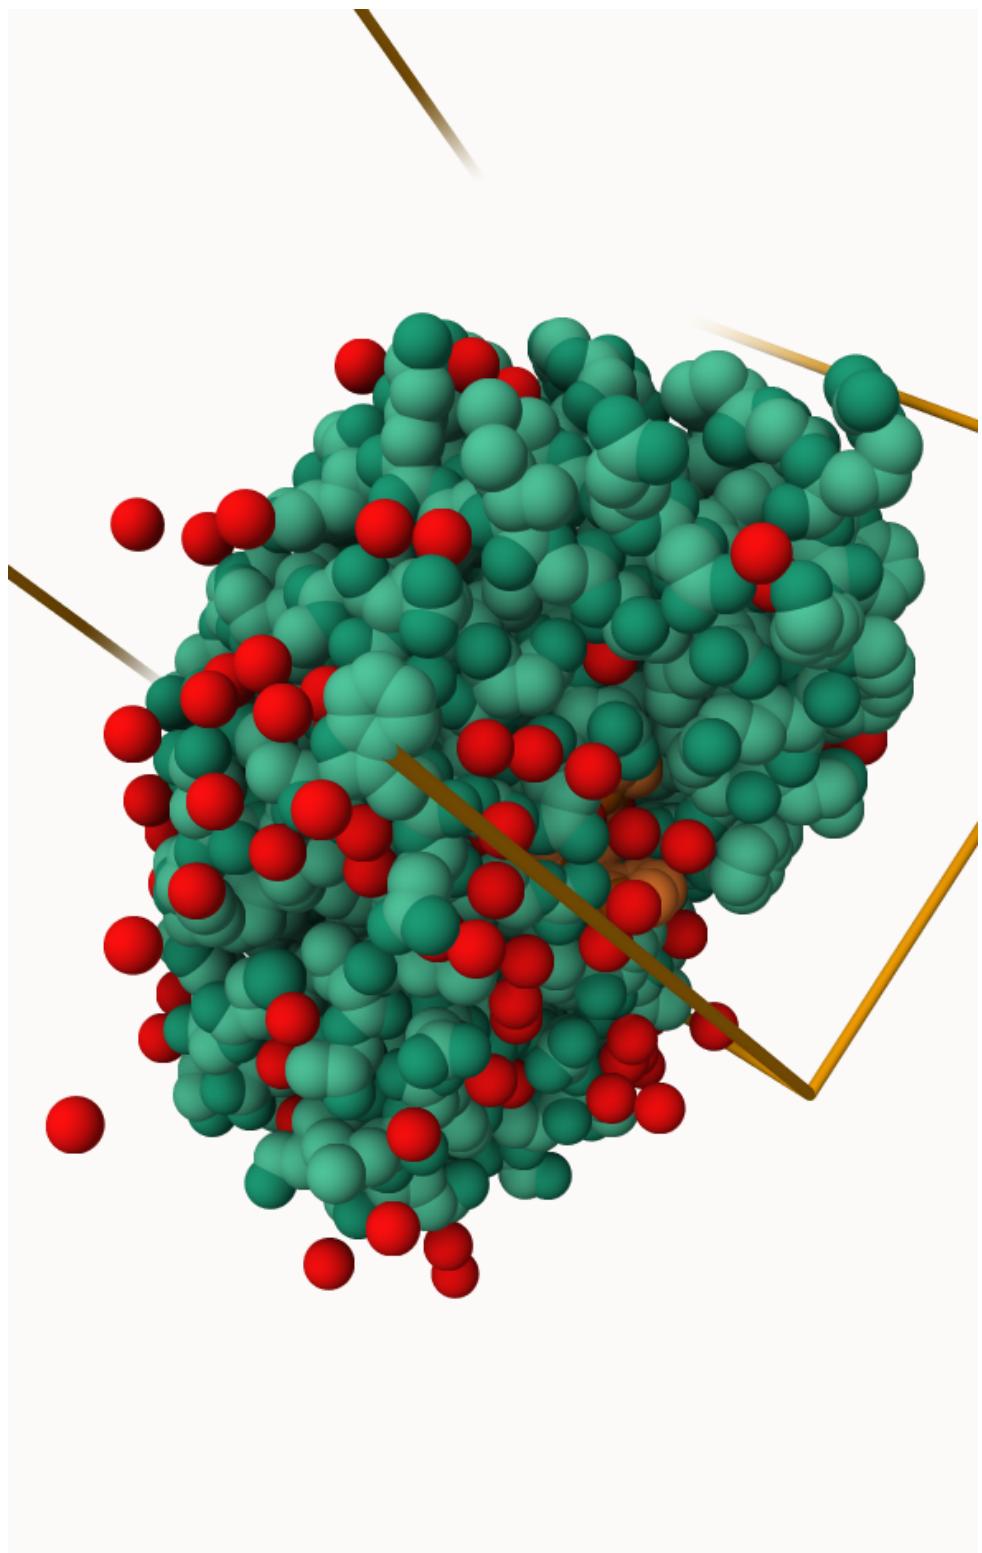


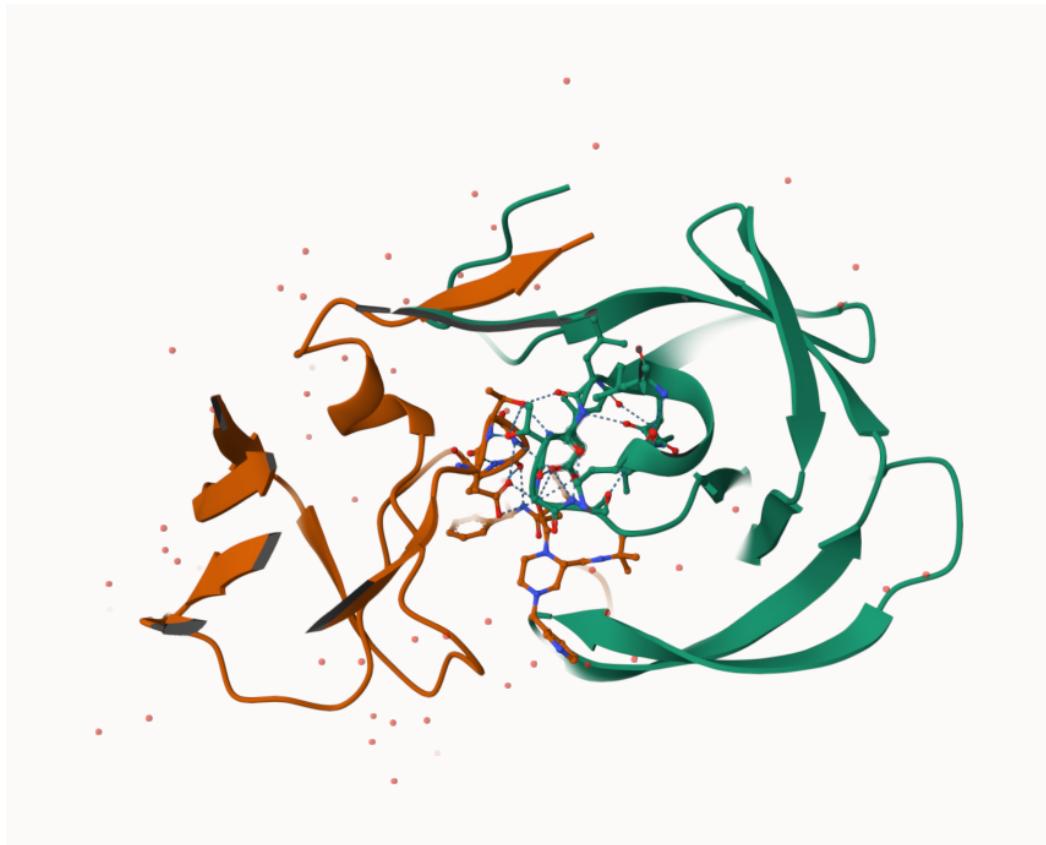
Figure 1: first view of HIV-Pr dimer with bound inhibitor

Q. Generate and Insert an image of the HIV-Pr cartoon colored by secondary structure, showing the inhibitor (ligand) in ball and stick.



Figure 2: first view of HIV-Pr showing the inhibitor in ball and stick

Q. One final image showing catalytic APS 25 as ball and stick and the all-important



activities

Q4. Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure?

Water ( $\text{H}_2\text{O}$ ) has three atoms, but in X-ray crystallography the hydrogen atoms barely scatter X-rays because they have very little electrons. Only the oxygen atom has enough electrons to be easily identified by the X-ray crystallography. Hydrogen atoms have very few electrons so it contributes almost no signal so they are usually invisible in electron-density maps. Meanwhile, the oxygen atom is detectable since it is more electron-dense. Therefore, this is why we just see one atom per water molecule in this structure.

Q5. There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have?

Yes I was able to identify the critical “conserved” water molecule in the binding site. The water molecule is HOH 326 and it most likely participates in catalysis.

Q6. Generate and save a figure clearly showing the two distinct chains of HIV-protease along with the ligand. You might also consider showing the catalytic residues ASP 25 in each chain and the critical water (we recommend “Ball & Stick” for these side-chains). Add this figure to your Quarto document.

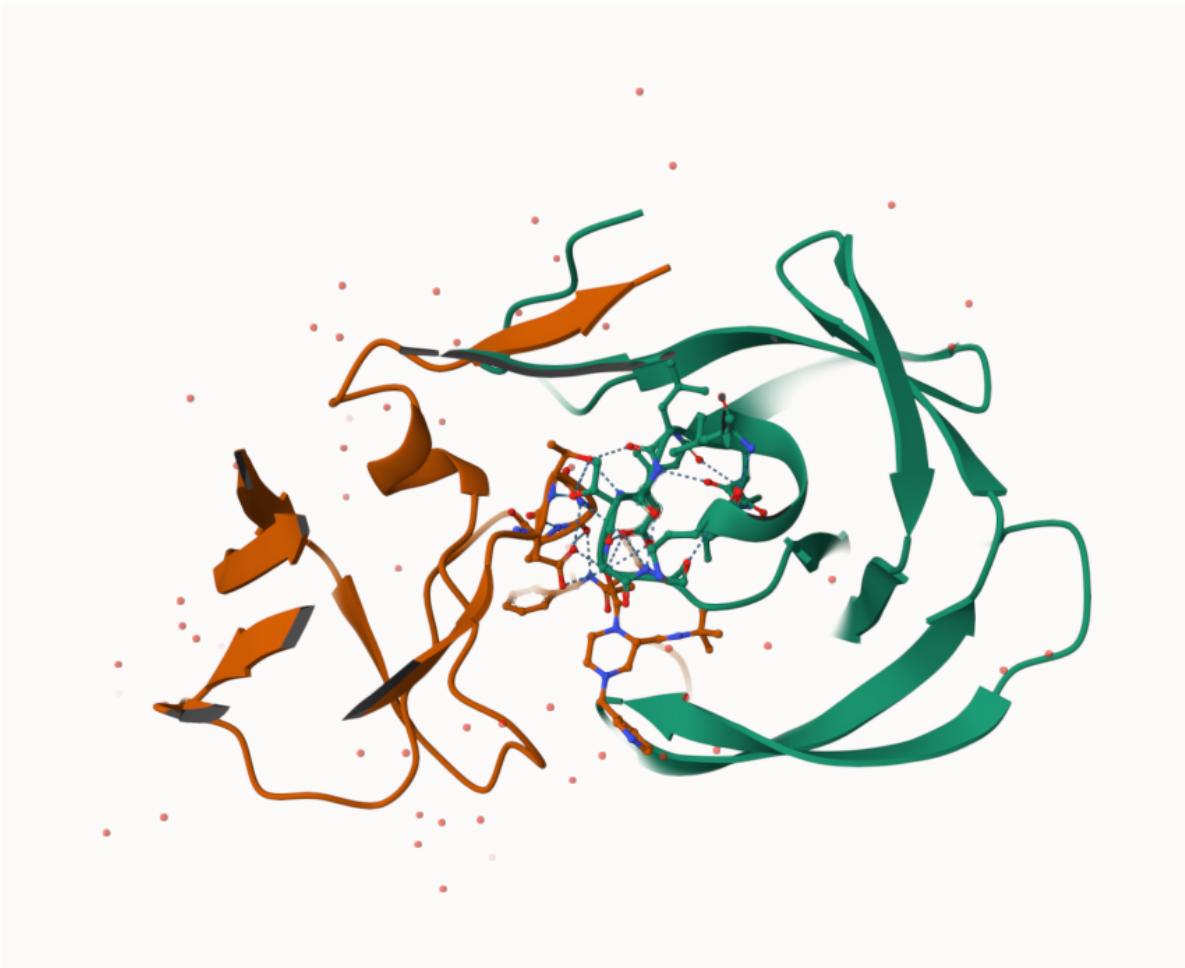


Figure 3: first view of HIV-Pr along with ligand

Discussion Topic: can you think of a way in which indinavir, or even larger ligands and substrates could enter the binding site?

The indinavir and other larger ligands can enter the binding site through transient opening of the flexible flap regions of the HIV-1 protease which provides access to the active site. So the ligand enters when the flaps adopt an open conformation, the ligand diffuses into the active site, and then the flaps close over the ligand, trapping it in place. This is known as an induced fit.

```
library(bio3d)
```

```
Warning: package 'bio3d' was built under R version 4.4.3
```

```
png1 <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
png1
```

```
Call: read.pdb(file = "1hsg")
```

Total Models#: 1

Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)

Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172 (residues: 128)

Non-protein/nucleic resid values: [ HOH (127), MK1 (1) ]

Protein sequence:

```
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPCKMIGGIGGFVKVRQYD  
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE  
ALLDTGADDTVLEEMSLPGRWPCKMIGGIGGFVKVRQYDQILIEICGHKAIGTVLVGPTP  
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,  
       calpha, remark, call
```

```
head(png1$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	29.361	39.686	5.862	1	38.10
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	30.307	38.663	5.319	1	40.62
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	29.760	38.071	4.022	1	42.64
4	ATOM	4	O	<NA>	PRO	A	1	<NA>	28.600	38.302	3.676	1	43.40
5	ATOM	5	CB	<NA>	PRO	A	1	<NA>	30.508	37.541	6.342	1	37.87
6	ATOM	6	CG	<NA>	PRO	A	1	<NA>	29.296	37.591	7.162	1	38.40
	segid	elesy	charge										
1	<NA>	N	<NA>										
2	<NA>	C	<NA>										
3	<NA>	C	<NA>										

```

4 <NA>      0 <NA>
5 <NA>      C <NA>
6 <NA>      C <NA>

```

`pdbsq(png1)`

```

1   2   3   4   5   6   7   8   9   10  11  12  13  14  15  16  17  18  19  20
"P" "Q" "I" "T" "L" "W" "Q" "R" "P" "L" "V" "T" "I" "K" "I" "G" "G" "Q" "L" "K"
21  22  23  24  25  26  27  28  29  30  31  32  33  34  35  36  37  38  39  40
"E" "A" "L" "L" "D" "T" "G" "A" "D" "D" "T" "V" "L" "E" "E" "M" "S" "L" "P" "G"
41  42  43  44  45  46  47  48  49  50  51  52  53  54  55  56  57  58  59  60
"R" "W" "K" "P" "K" "M" "I" "G" "G" "I" "G" "F" "I" "K" "V" "R" "Q" "Y" "D"
61  62  63  64  65  66  67  68  69  70  71  72  73  74  75  76  77  78  79  80
"Q" "I" "L" "I" "E" "I" "C" "G" "H" "K" "A" "I" "G" "T" "V" "L" "V" "G" "P" "T"
81  82  83  84  85  86  87  88  89  90  91  92  93  94  95  96  97  98  99  1
"P" "V" "N" "I" "I" "G" "R" "N" "L" "L" "T" "Q" "I" "G" "C" "T" "L" "N" "F" "P"
2   3   4   5   6   7   8   9   10  11  12  13  14  15  16  17  18  19  20  21
"Q" "I" "T" "L" "W" "Q" "R" "P" "L" "V" "T" "I" "K" "I" "G" "G" "Q" "L" "K" "E"
22  23  24  25  26  27  28  29  30  31  32  33  34  35  36  37  38  39  40  41
"A" "L" "L" "D" "T" "G" "A" "D" "D" "T" "V" "L" "E" "E" "M" "S" "L" "P" "G" "R"
42  43  44  45  46  47  48  49  50  51  52  53  54  55  56  57  58  59  60  61
"W" "K" "P" "K" "M" "I" "G" "G" "I" "G" "F" "I" "K" "V" "R" "Q" "Y" "D" "Q"
62  63  64  65  66  67  68  69  70  71  72  73  74  75  76  77  78  79  80  81
"I" "L" "I" "E" "I" "C" "G" "H" "K" "A" "I" "G" "T" "V" "L" "V" "G" "P" "T" "P"
82  83  84  85  86  87  88  89  90  91  92  93  94  95  96  97  98  99
"V" "N" "I" "I" "G" "R" "N" "L" "L" "T" "Q" "I" "G" "C" "T" "L" "N" "F"

```

Q7. How many amino acid residues are there in this pdb object?

```

# Counts amino-acid residues across all chains in a PDB object
aa_filter <- png1$atom$resid %in% c(
  "ALA", "ARG", "ASN", "ASP", "CYS",
  "GLU", "GLN", "GLY", "HIS", "ILE",
  "LEU", "LYS", "MET", "PHE", "PRO",
  "SER", "THR", "TRP", "TYR", "VAL"
)
pairs <- paste(
  png1$atom$chain[aa_filter],
  png1$atom$resno[aa_filter],
  sep = "_"
)
length(unique(pairs))

```

```
[1] 198
```

Q8. Name one of the two non-protein residues?

HOH is one of the two non-protein residue.

```
unique(png1$atom$resid)
```

```
[1] "PRO" "GLN" "ILE" "THR" "LEU" "TRP" "ARG" "VAL" "LYS" "GLY" "GLU" "ALA"  
[13] "ASP" "MET" "SER" "PHE" "TYR" "CYS" "HIS" "ASN" "MK1" "HOH"
```

Q9. How many protein chains are in this structure?

There are 2 protein chains in this structure.

```
aa <- c("ALA", "ARG", "ASN", "ASP", "CYS",  
       "GLU", "GLN", "GLY", "HIS", "ILE",  
       "LEU", "LYS", "MET", "PHE", "PRO",  
       "SER", "THR", "TRP", "TYR", "VAL")  
  
aa_filter <- png1$atom$resid %in% aa  
  
protein_chains <- unique(png1$atom$chain[aa_filter])  
length(protein_chains)
```

```
[1] 2
```

### **Install – Install packages in the R console NOT your Rmd/Quarto file**

```
install.packages("bio3d") install.packages("NGLVieweR")  
install.packages("remotes") remotes::install_github("bioboot/bio3dview")  
install.packages("BiocManager") BiocManager::install("msa")
```

Q10. Which of the packages above is found only on BioConductor and not CRAN?

msa is found onnly on BioConductor and not CRAN.

Q11. Which of the above packages is not found on BioConductor or CRAN?

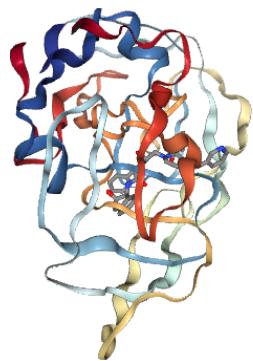
bio3dview is not found on BioConductor or CRAN.

Q12. True or False? Functions from the `pak` package can be used to install packages from GitHub and BitBucket?

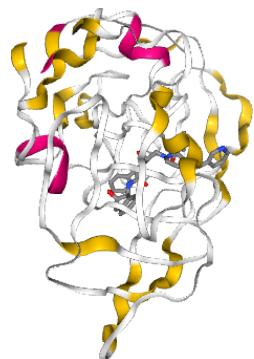
TRUE

Let's try out the new **bio3dview** package that is not yet on CRAN. We can use the **remotes** package to install any R package from GitHub.

```
library(bio3dview)  
  
view.pdb(png1)
```



```
view.pdb(png1, colorScheme = "sse")
```

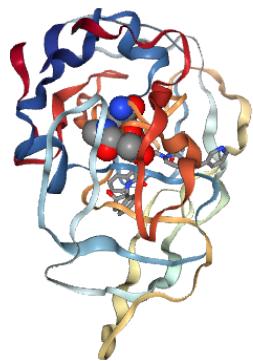


## Quick viewing of PDBs

```
library(bio3dview)

sele <- atom.select(png1, resno=25)

view.pdb(png1, backgroundColor = "pink",
         highlight = sele,
         highlight.style = "spacefill")
```



## Prediction of Protein Flexibility

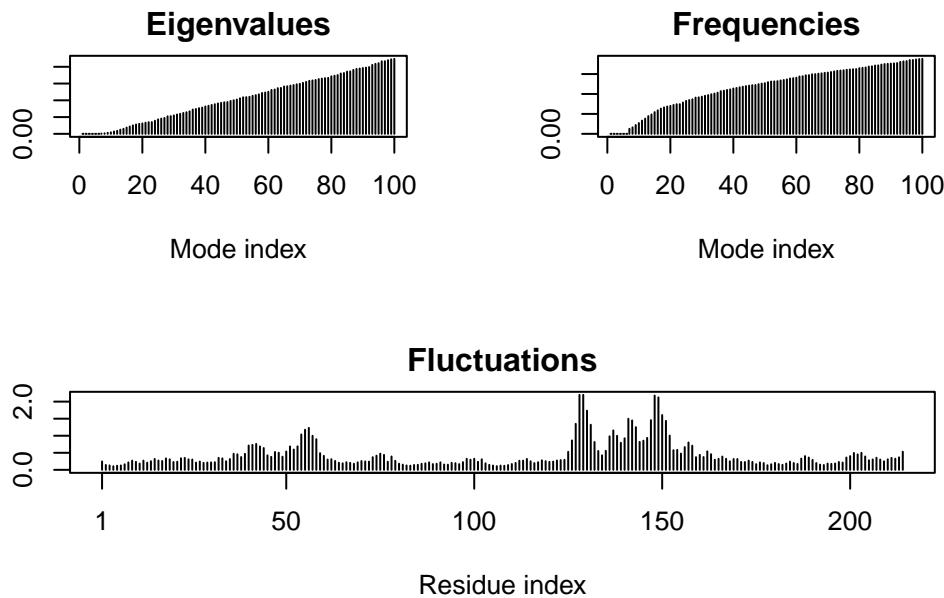
```
adk <- read.pdb("6s36")
```

Note: Accessing on-line PDB file  
PDB has ALT records, taking A only, rm.alt=TRUE

```
m <- nma(adk)
```

Building Hessian... Done in 0.01 seconds.  
Diagonalizing Hessian... Done in 0.14 seconds.

```
plot(m)
```



Write out our results as a wee trajectory movie:

```
mktrj(m, file="results.pdb")
view.nma(m)
```

```
library(bio3d)  
  
aa <- get.seq("1ake_A")
```

Warning in get.seq("1ake\_A"): Removing existing file: seqs.fasta

Fetching... Please wait. Done.

```
aa
```

1 pdb 1AKE A	MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRAAVKSGSELGKQAKDIMDAGKLVT 1 1	60 60
61 pdb 1AKE A	DELVIALVKERIAQEDCRNGFLLDGFPRTIPQADAMKEAGINVDYVLEFDVPDELIVDRI 61 61	120 120
121 pdb 1AKE A	VGRRVHAPSGRKYHVKFNPPKVEGKDDVTGEELTRKDDQEETVRKRLVEYHQMTAPLIG 121 121	180 180
181 pdb 1AKE A	YYSKAEAGNTKYAKVDGTPVAEVRADLEKILG 181 181	214 214

Call:

```
read.fasta(file = outfile)
```

Class:

```
fasta
```

Alignment dimensions:

```
1 sequence rows; 214 position columns (214 non-gap, 0 gap)
```

```
+ attr: id, ali, call
```

Q13. How many amino acids are in this sequence, i.e. how long is this sequence?

There are 214 amino acids in this sequence.

```
length(aa$ali)
```

```
[1] 214
```